

GPT-2 Tutorial: Understanding GPT from Scratch

A complete guide to understanding every component of a GPT (Generative Pre-trained Transformer) language model, explained as if you're learning it for the first time.

Author: Kevin Thomas (ket189@pitt.edu)

License: MIT

Note: This tutorial uses simplified, educational variable names (like `embedding_dim`) for clarity. The production code in `micro_gpt.py` follows OpenAI GPT-2 naming conventions (`n_embd`, `n_head`, `n_layer`).

Table of Contents

1. Introduction: What is a Language Model?
 2. Tokenization: Breaking Text into Pieces
 3. Token Embeddings: Numbers to Vectors
 4. Positional Embeddings: Teaching Position
 5. The Residual Stream: The Information Highway
 6. Self-Attention: Learning What's Important
 7. Multi-Head Attention: Multiple Perspectives
 8. Feed-Forward Networks: Processing Information
 9. Transformer Blocks: Putting It Together
 10. The Complete GPT Architecture
 11. Training: Teaching the Model
 12. Generation: Creating New Text
 13. Putting It All Together
-

1. Introduction: What is a Language Model?

The Big Picture

Imagine you're playing a game where someone gives you the beginning of a sentence, and you have to guess the next word. For example:

- “The cat sat on the ____” → You’d probably guess “mat” or “chair”
- “I love eating ____” → You might guess “pizza” or “ice cream”

That’s exactly what a language model does! It predicts what word (or token) comes next in a sequence.

Why This Matters

Language models power: - ChatGPT and similar chatbots - Auto-complete in your phone - Translation services - Code completion in IDEs

How GPT is Different

GPT (Generative Pre-trained Transformer) uses a special architecture called a **transformer** that can: 1. Look at all previous words at once (not just the last few) 2. Understand which words are important to pay attention to 3. Generate coherent, human-like text

Let's break down every single piece!

2. Tokenization: Breaking Text into Pieces

What is Tokenization?

Before a computer can understand text, we need to convert words into numbers. But we don't just convert whole words - we break text into smaller pieces called **tokens**.

Example

The sentence "Hello, world!" might be tokenized as:

```
"Hello, world!" → ["Hello", ",", "world", "!"] → [15496, 11, 995, 0]
```

Each token gets a unique number (an ID) from a vocabulary.

Why Not Just Use Words?

Tokens are better than words because: - **Handles unknown words**: "unhappiness" = "un" + "happiness" - **Smaller vocabulary**: We only need ~50,000 tokens instead of millions of words - **Works for any language**: Even handles code and symbols!

In Our Code

```
# We use tiktoken (OpenAI's tokenizer)
encoding = tiktoken.get_encoding("gpt2")
tokens = encoding.encode("Hello, world!") # [15496, 11, 995, 0]
```

Key Point: After tokenization, our text is now a sequence of numbers (token IDs).

3. Token Embeddings: Numbers to Vectors

The Problem

We have token IDs like [15496, 11, 995], but neural networks can't learn much from single numbers. We need **vectors** (lists of numbers) that can capture meaning.

What is an Embedding?

An embedding converts each token ID into a vector (a list of many numbers). Think of it like this:

Before (Token ID):

```
"cat" → 145
```

After (Embedding Vector):

```
"cat" → [0.2, -0.5, 0.8, 0.1, -0.3, ...] (maybe 64 or 384 numbers!)
```

Why Vectors?

Vectors can capture relationships: - “king” - “man” + “woman” = “queen” - Similar words have similar vectors - The model learns these relationships during training!

The Analogy

Think of embeddings like describing a person: - Token ID: Just a name “John” (one number) - Embedding: Height, weight, age, personality traits, interests... (many numbers)

The more numbers, the more detailed the description!

In Our Code

```
class GPT2(nn.Module):
    def __init__(self, vocab_size=50257, embedding_dim=384, ...):
        # Create embedding table: vocab_size × embedding_dim
        self.token_embedding = nn.Embedding(vocab_size, embedding_dim)
```

If `vocab_size=50257` and `embedding_dim=384`: - We have 50,257 tokens - Each token becomes a 384-dimensional vector - Total: $50,257 \times 384 = 19$ million learnable numbers!

Example

```
# Token IDs (integers)
tokens = torch.tensor([145, 892, 33]) # Shape: (3,)

# After embedding (vectors)
embedded = token_embedding(tokens) # Shape: (3, 384)
# Each token is now a 384-dimensional vector!
```

Key Point: Embeddings transform token IDs into rich, meaningful vectors that capture semantic information.

4. Positional Embeddings: Teaching Position

The Problem

Look at these two sentences: 1. “The cat chased the mouse” 2. “The mouse chased the cat”

Same words, **completely different meanings!** The model needs to know **where** each word appears in the sequence.

The Solution: Positional Embeddings

Just like we embed tokens, we also embed positions:

```

Position 0 → [0.5, 0.2, -0.1, ...]
Position 1 → [0.3, -0.4, 0.6, ...]
Position 2 → [0.1, 0.8, -0.2, ...]
...

```

How It Works

For the sequence “The cat sat”:

Step 1: Token Embeddings

```

"The" → token_embedding[145] → [0.2, -0.5, 0.8, ...]
"cat" → token_embedding[892] → [0.1, 0.3, -0.2, ...]
"sat" → token_embedding[33] → [0.4, 0.1, 0.5, ...]

```

Step 2: Position Embeddings

```

Position 0 → position_embedding[0] → [0.5, 0.2, -0.1, ...]
Position 1 → position_embedding[1] → [0.3, -0.4, 0.6, ...]
Position 2 → position_embedding[2] → [0.1, 0.8, -0.2, ...]

```

Step 3: Add Them Together!

```

"The" (pos 0) = token_emb[145] + pos_emb[0]
"cat" (pos 1) = token_emb[892] + pos_emb[1]
"sat" (pos 2) = token_emb[33] + pos_emb[2]

```

Why Addition?

We add (not concatenate) because: 1. Keeps the same dimension (384 stays 384) 2. The model learns to separate content from position during training 3. It’s simpler and works well!

In Our Code

```

def _embed_tokens(self, idx: torch.Tensor) -> torch.Tensor:
    T = idx.size(1) # Sequence length
    # Get token embeddings
    tok_emb = self.token_embedding(idx) # (B, T, embedding_dim)
    # Get position embeddings
    pos_emb = self.position_embedding(torch.arange(T)) # (T, embedding_dim)
    # Add them together!
    return tok_emb + pos_emb # (B, T, embedding_dim)

```

Example with Numbers

```

# Input: 2 sequences of 3 tokens each
idx = torch.tensor([[145, 892, 33],
                   [567, 234, 891]]) # Shape: (2, 3)

# After token embedding
tok_emb.shape # (2, 3, 384)

```

```

# After position embedding
pos_emb.shape # (3, 384) - same for all sequences!

# After adding
result.shape # (2, 3, 384)

```

Key Point: Positional embeddings give the model a sense of order, so it knows “the cat chased the mouse” is different from “the mouse chased the cat”.

5. The Residual Stream: The Information Highway

The Problem

Imagine you’re passing a note through a chain of 6 people. By the time it reaches person 6, the original message might be garbled or lost!

Neural networks have the same problem - information can get lost or corrupted as it passes through many layers.

The Solution: Residual Connections (Skip Connections)

Instead of just passing information forward, we **add a shortcut**:

Before (no residual):

Input → Layer 1 → Layer 2 → Layer 3 → Output

After (with residual):

Input → Layer 1 → + → Layer 2 → + → Layer 3 → Output
 ↑ ↑

How It Works

Instead of:

```
x = layer(x) # x is completely replaced
```

We do:

```
x = x + layer(x) # Add the layer output to the original x!
```

The Analogy

Think of it like editing a document: - **Without residual**: Rewrite the whole document from scratch each time - **With residual**: Start with the original and only make changes/additions

The original information always flows through!

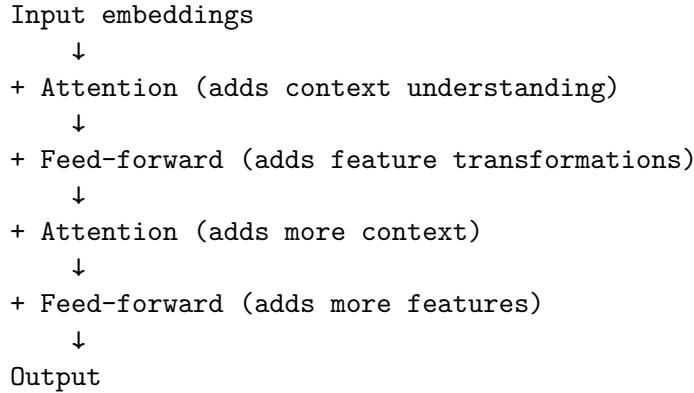
Why This is Powerful

1. **Gradients flow better**: Makes training deep networks easier
2. **Information preservation**: Original input never gets lost

3. Allows depth: We can stack 100+ layers successfully!

The Residual Stream Concept

In transformers, we think of the data flow as a **residual stream** - a river of information that flows through the network, with each layer adding to it:



In Our Code

```
class Block(nn.Module):
    def forward(self, x: torch.Tensor) -> torch.Tensor:
        # Residual connection around attention
        x = x + self.sa(self.ln1(x))  # x = x + attention_output

        # Residual connection around feed-forward
        x = x + self.ffwd(self.ln2(x))  # x = x + feedforward_output

    return x
```

Example Flow

```
# Start with embeddings
x = embeddings  # Shape: (2, 10, 384)

# Block 1
x = x + attention(x)      # Add attention's contribution
x = x + feedforward(x)    # Add feedforward's contribution

# Block 2
x = x + attention(x)      # Add more attention
x = x + feedforward(x)    # Add more processing

# Original embeddings are still "in there" plus all the additions!
```

Key Point: Residual connections create an “information superhighway” that preserves important information while allowing the network to learn complex transformations.

6. Self-Attention: Learning What's Important

The Big Idea

When you read a sentence, you naturally pay attention to certain words more than others. Attention mechanisms let the model do the same!

The Motivation

Consider: “The animal didn’t cross the street because **it** was too tired”

What does “it” refer to? The animal or the street? You know it’s the animal because you pay attention to context clues like “tired”.

Self-attention lets the model learn these relationships!

The Three Players: Query, Key, Value

Think of attention like a search engine:

1. **Query (Q)**: “What am I looking for?”
 - Like typing a search query: “tired animal”
2. **Key (K)**: “What do I have?”
 - Like document titles: “street description”, “animal description”
3. **Value (V)**: “What information can I provide?”
 - The actual content of the documents

The Attention Mechanism (Step by Step)

Step 1: Create Q, K, V For each word, we create three vectors:

```
# For each token embedding (384-dimensional)
q = query_layer(x)    # "What should I look for?"
k = key_layer(x)      # "What do I represent?"
v = value_layer(x)    # "What information do I carry?"
```

Step 2: Calculate Attention Scores We compare each Query with all Keys to see how relevant they are:

```
# Dot product: how similar is each query to each key?
scores = q @ k.transpose(-2, -1) # (B, T, T)
```

Example: For “it” in our sentence:

```
"it" queries:
- "animal" key: HIGH similarity (0.9)
- "street" key: LOW similarity (0.2)
- "tired" key: MEDIUM similarity (0.6)
```

Step 3: Scale the Scores Divide by $\sqrt{\text{head_size}}$ to keep gradients stable:

```
scores = scores / (head_size ** 0.5)
```

Why? Prevents scores from getting too large (which would make softmax produce near-0 or near-1 values).

Step 4: Causal Masking (for GPT) **Critical for language models!** We can't look at future words:

```
# Mask out future positions (set them to -infinity)
scores = scores.masked_fill(mask == 0, float('-inf'))
```

Example:

Original scores:	After masking:
[0.5 0.3 0.8]	[0.5 -inf -inf] ← position 0 can't see 1,2
[0.4 0.6 0.2]	→ [0.4 0.6 -inf] ← position 1 can't see 2
[0.7 0.5 0.9]	[0.7 0.5 0.9] ← position 2 can see all

This ensures: **We only predict based on past context, never future!**

Step 5: Softmax (Make it a Probability) Convert scores to probabilities that sum to 1:

```
attention_weights = softmax(scores) # (B, T, T)
```

Example for "it":

Before softmax:	After softmax:
animal: 0.9 →	animal: 0.55 (55% attention)
street: 0.2 →	street: 0.12 (12% attention)
tired: 0.6 →	tired: 0.33 (33% attention)
	Total: 1.00 (100%)

Step 6: Weighted Sum Use attention weights to mix the Values:

```
output = attention_weights @ v # (B, T, head_size)
```

What this means: - "it" becomes 55% animal + 12% street + 33% tired - The output captures context from all related words!

Complete Example

```
# Sentence: "The cat sat"
# Positions: 0:"The", 1:"cat", 2:"sat"

# For position 2 ("sat"), the attention might look like:
Attention weights for "sat":
↓
0.1 ← "The" (10% attention - not very relevant)
0.6 ← "cat" (60% attention - the subject!)
0.3 ← "sat" (30% attention - the word itself)

# Output for "sat" becomes:
output = 0.1 * value["The"] + 0.6 * value["cat"] + 0.3 * value["sat"]
```

In Our Code

```
class SelfAttentionHead(nn.Module):
    def forward(self, x: torch.Tensor) -> torch.Tensor:
        # Step 1: Create Q, K, V
        k = self.key(x)      # (B, T, head_size)
        q = self.query(x)    # (B, T, head_size)
        v = self.value(x)   # (B, T, head_size)

        # Steps 2-5: Compute attention weights
        wei = q @ k.transpose(-2, -1)  # (B, T, T) - Step 2
        wei = wei / (k.size(-1) ** 0.5) # Step 3: scale
        wei = wei.masked_fill(self.tril[:T, :T] == 0, float('-inf')) # Step 4: mask
        wei = F.softmax(wei, dim=-1)  # Step 5: normalize

        # Step 6: Weighted sum
        output = wei @ v # (B, T, head_size)
    return output
```

Visualization

Input: [The, cat, sat]

After attention, "sat" has looked at its context:
"sat" = 10% "The" + 60% "cat" + 30% "sat" (itself)

So "sat" now knows it's related to a "cat" doing the action!

Key Point: Self-attention lets each word look at all other words (in the past) and decide which ones are important for understanding its meaning in context.

7. Multi-Head Attention: Multiple Perspectives

The Problem with Single Attention

One attention head can only learn one type of relationship. But language has many types of relationships:

- **Syntactic:** “What’s the verb?” “What’s the subject?”
- **Semantic:** “What does this refer to?” “What’s related?”
- **Positional:** “What came recently?” “What’s nearby?”

The Solution: Multiple Heads

Instead of one attention mechanism, use many **in parallel!**

Think of it like asking multiple experts: - Expert 1 (Head 1): Focuses on grammar relationships - Expert 2 (Head 2): Focuses on semantic meaning - Expert 3 (Head 3): Focuses on nearby words - Expert 4 (Head 4): Focuses on sentence structure

Each head learns to pay attention to different things!

How It Works

Step 1: Create Multiple Heads Instead of one attention head with `head_size=384`, we create multiple smaller heads:

```
# Single head: 384 dimensions
# 6 heads: 6 × 64 = 384 dimensions total

heads = [
    SelfAttentionHead(embedding_dim=384, head_size=64), # Head 1
    SelfAttentionHead(embedding_dim=384, head_size=64), # Head 2
    SelfAttentionHead(embedding_dim=384, head_size=64), # Head 3
    # ... 3 more heads
]
```

Step 2: Run All Heads in Parallel Each head processes the input independently:

```
head_outputs = [head(x) for head in heads]
# Each output: (B, T, 64)
```

Step 3: Concatenate Join all head outputs side-by-side:

```
output = torch.cat(head_outputs, dim=-1) # (B, T, 384)
# 64 + 64 + 64 + 64 + 64 + 64 = 384
```

Step 4: Project Mix the head outputs together:

```
output = projection_layer(output) # (B, T, 384)
```

Example

Sentence: "The cat sat on the mat"

Head 1 might learn:

"sat" pays attention to "cat" (subject-verb relationship)

Head 2 might learn:

"sat" pays attention to "on" (verb-preposition relationship)

Head 3 might learn:

"sat" pays attention to nearby words (local context)

Head 4 might learn:

"sat" pays attention to the start of sentence (sentence structure)

Combined output:

"sat" has information about its subject, object, position, and context!

Why Smaller Heads?

Instead of one 384-dim head, we use 6×64 -dim heads:

Advantages: 1. **Specialization:** Each head learns different patterns 2. **Same computation:** $6 \times 64 = 384$ (same total dimensions) 3. **More flexible:** Can learn multiple relationships simultaneously

The Math:

Option 1: 1 head \times 384 dimensions = 384 total

Option 2: 6 heads \times 64 dimensions = 384 total (same!)

But Option 2 can learn 6 different attention patterns!

In Our Code

```
class MultiHeadAttention(nn.Module):
    def __init__(self, embedding_dim, num_heads=6):
        head_size = embedding_dim // num_heads # 384 // 6 = 64

        # Create multiple heads
        self.heads = nn.ModuleList([
            SelfAttentionHead(embedding_dim, head_size)
            for _ in range(num_heads)
        ])

        # Projection layer to mix head outputs
        self.proj = nn.Linear(embedding_dim, embedding_dim)

    def forward(self, x):
        # Run all heads in parallel
        out = torch.cat([h(x) for h in self.heads], dim=-1)
        # Mix them together
        return self.proj(out)
```

Visualization

Input: (B, T, 384)

↓

↓	↓	↓	↓	↓	↓
Head1	Head2	Head3	Head4	Head5	Head6
(64)	(64)	(64)	(64)	(64)	(64)

↓

Concatenate

↓

(B, T, 384)

```
↓  
Projection  
↓  
Output: (B, T, 384)
```

Real-World Analogy

Imagine writing a book review. You might consider: - Plot (Head 1) - Characters (Head 2) - Writing style (Head 3) - Pacing (Head 4) - Themes (Head 5) - Overall impression (Head 6)

Each “head” focuses on one aspect, then you combine all perspectives for your final review!

Key Point: Multi-head attention lets the model look at sequences from multiple perspectives simultaneously, learning different types of relationships in parallel.

8. Feed-Forward Networks: Processing Information

What Comes After Attention?

Attention helps words understand their context by looking at other words. But we also need to **process** and **transform** that information!

That’s where Feed-Forward Networks (FFN) come in.

What is a Feed-Forward Network?

It’s a simple two-layer neural network that processes each position independently:

$$\text{FFN}(x) = \text{ReLU}(x @ W_1 + b_1) @ W_2 + b_2$$

Think of it as: **expand** → **transform** → **compress**

The Architecture

```
Input: 384 dimensions  
↓  
Expand: 384 → 1536 dimensions (4× larger!)  
↓  
ReLU activation (non-linearity)  
↓  
Compress: 1536 → 384 dimensions (back to original)  
↓  
Output: 384 dimensions
```

Why Expand Then Compress?

The Analogy: Imagine explaining a concept: 1. **Expand:** Break it down into detailed examples (more information) 2. **Process:** Think about each detail, make connections 3. **Compress:** Synthesize back to a clear, refined explanation

In the network: 1. **Expand (384→1536):** Create more “features” to work with 2. **ReLU:** Add non-linearity (enables learning complex patterns) 3. **Compress (1536→384):** Distill the useful information

The $4\times$ Rule

Transformers typically use a $4\times$ expansion ratio: - Input: 384 dimensions - Hidden: $384 \times 4 = 1536$ dimensions - Output: 384 dimensions

Why $4\times$? - More capacity for learning complex transformations - But not so large that training becomes too slow - It’s a sweet spot found through experimentation!

What Does the FFN Learn?

While attention learns **relationships between positions**, FFN learns **transformations at each position**:

Attention: "sat" learns it's related to "cat"

FFN: Transforms "sat related to cat" → "past-tense verb action by feline"

FFN can learn patterns like: - “Plural nouns need plural verbs” - “Past tense patterns” - “Semantic categories” (animals, food, actions) - “Syntactic structures”

Position-wise Processing

Important: FFN processes each position independently:

```
# These are equivalent:
output = ffn(x)  # Process whole sequence

# Same as:
output = torch.stack([ffn(x[:, i, :]) for i in range(T)], dim=1)
# Process each position separately!
```

Why independent? - After attention mixed information across positions - Now we refine each position’s representation - Attention = communication, FFN = computation

In Our Code

```
class FeedForward(nn.Module):
    def __init__(self, n_embd, dropout=0.1):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(n_embd, 4 * n_embd),      # Expand: 384 → 1536
            nn.ReLU(),                          # Non-linearity
            nn.Dropout(dropout),               # Regularization
            nn.Linear(4 * n_embd, n_embd),     # Compress: 1536 → 384
            nn.Dropout(dropout),               # More regularization
        )

    def forward(self, x):
        return self.net(x)
```

Step-by-Step Example

```
# Input from attention
x = torch.randn(2, 10, 384) # (batch=2, seq_len=10, dim=384)

# Step 1: Expand
hidden = linear1(x) # (2, 10, 1536) - 4x bigger!

# Step 2: ReLU activation
hidden = relu(hidden) # (2, 10, 1536) - now non-linear

# Step 3: Compress
output = linear2(hidden) # (2, 10, 384) - back to original size

# Each of the 10 positions processed independently!
```

Why ReLU?

ReLU (Rectified Linear Unit) adds non-linearity:

$$\text{ReLU}(x) = \max(0, x)$$

```
# Examples:
ReLU(-0.5) = 0 # Negative → 0
ReLU(0.3) = 0.3 # Positive → unchanged
ReLU(-2.1) = 0 # Negative → 0
ReLU(1.7) = 1.7 # Positive → unchanged
```

Without non-linearity, the network would just be a linear transformation (boring!). ReLU lets it learn complex patterns.

The Role in the Transformer Block

```
Input
  ↓
Attention (look at other positions)
  ↓
+ Add & Norm (residual connection)
  ↓
Feed-Forward (process each position)
  ↓
+ Add & Norm (residual connection)
  ↓
Output
```

Attention = “communicate” (positions talk to each other) FFN = “compute” (each position thinks independently)

Real-World Analogy

Think of a team project: 1. **Attention (Meeting)**: Everyone shares their ideas and listens to others 2. **FFN (Individual Work)**: Each person goes away and processes what they learned

Attention = collaboration, FFN = individual processing

Key Point: Feed-forward networks provide the computational power to transform and refine the information gathered by attention, processing each position independently.

9. Transformer Blocks: Putting It Together

The Complete Building Block

A **Transformer Block** combines everything we've learned so far into one powerful unit:

```
Input
  ↓
Layer Norm
  ↓
Multi-Head Attention
  ↓
+ (Residual) ←
  ↓
Layer Norm
  ↓
Feed-Forward
  ↓
+ (Residual) ←
  ↓
Output
```

The Six Key Components

1. Layer Normalization (Pre-Norm) Before each sub-layer, we normalize:

```
# Layer norm: mean=0, std=1 for each position
```

```
x_norm = (x - mean) / std
```

```
# Example:
```

```
x = [1, 2, 3, 4, 5]
```

```
x_norm = [-1.4, -0.7, 0, 0.7, 1.4] # mean=0, std=1
```

Why? - Stabilizes training (prevents exploding/vanishing gradients) - Each layer gets consistent input ranges - Makes training faster and more reliable

Pre-Norm vs Post-Norm: We do it **before** the sub-layer (pre-norm), which works better for deep networks.

2. Multi-Head Attention We covered this! Each position looks at all previous positions.

3. First Residual Connection

```
x = x + attention(layer_norm(x))
```

The input “skips around” attention and gets added back!

4. Layer Normalization (Again) Normalize before the feed-forward layer.

5. Feed-Forward Network We covered this! Transform each position independently.

6. Second Residual Connection

```
x = x + feed_forward(layer_norm(x))
```

Again, the input skips around and gets added back!

Why This Order?

Pre-Norm Architecture (what we use):

```
x = x + attention(layer_norm(x))
x = x + feed_forward(layer_norm(x))
```

Advantages: - Easier to train deep networks - More stable gradients - Better performance with many layers (6, 12, 24+ blocks)

In Our Code

```
class Block(nn.Module):
    def __init__(self, embedding_dim, block_size, n_heads, dropout=0.1):
        super().__init__()
        # Components
        self.sa = MultiHeadAttention(embedding_dim, block_size, n_heads, dropout)
        self.ffwd = FeedForward(embedding_dim, dropout)
        self.ln1 = nn.LayerNorm(embedding_dim)  # Before attention
        self.ln2 = nn.LayerNorm(embedding_dim)  # Before feed-forward

    def forward(self, x):
        # Attention block with residual
        x = x + self.sa(self.ln1(x))
        # Feed-forward block with residual
        x = x + self.ffwd(self.ln2(x))
        return x
```

Data Flow Example

```
# Start with embeddings
x = embeddings # (2, 10, 384)

# ===== Block 1 =====
# Step 1: Normalize
```

```

x_norm = layer_norm1(x) # (2, 10, 384)

# Step 2: Attention
attn_out = multi_head_attention(x_norm) # (2, 10, 384)

# Step 3: Add (residual)
x = x + attn_out # Original x + attention output

# Step 4: Normalize
x_norm = layer_norm2(x) # (2, 10, 384)

# Step 5: Feed-forward
ff_out = feed_forward(x_norm) # (2, 10, 384)

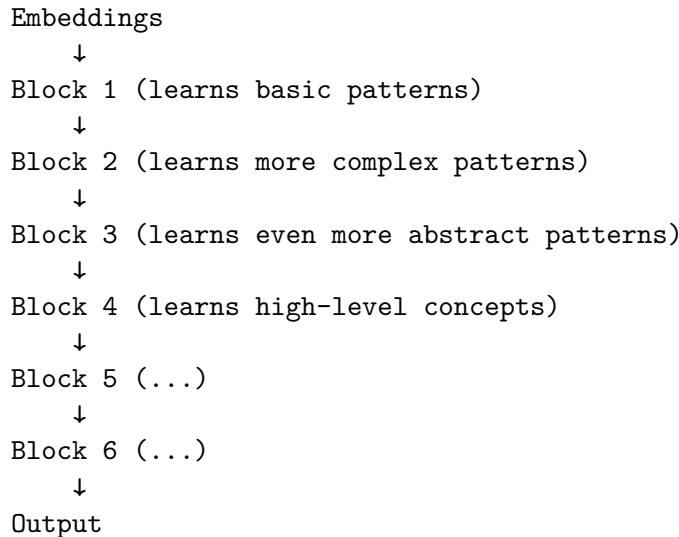
# Step 6: Add (residual)
x = x + ff_out # Previous x + feed-forward output

# After block 1: x now contains:
# - Original embeddings (via residuals)
# - Attention patterns learned
# - Feed-forward transformations

```

Stacking Blocks

The real power comes from stacking multiple blocks:



Each block can focus on different levels of abstraction: - Early blocks: syntax, simple patterns - Middle blocks: semantic relationships - Later blocks: high-level reasoning

Why Multiple Blocks Work

The Residual Stream Concept (revisited):

Think of the data flowing through as a stream, with each block adding layers of understanding:

```

Start: "cat sat mat"
↓
Block 1: "cat sat mat" + [syntax: subject-verb-object]
↓
Block 2: "cat sat mat" + [syntax] + [semantics: animal, action, object]
↓
Block 3: "cat sat mat" + [syntax] + [semantics] + [context: past tense scene]
↓
Block 4: "cat sat mat" + [syntax] + [semantics] + [context] + [inference: casual scene]
↓
...

```

Each block adds a “layer” of understanding while preserving all previous information!

Real-World Analogy

Think of reading comprehension: - **Block 1:** Identify words and grammar - **Block 2:** Understand sentence structure - **Block 3:** Grasp paragraph meaning - **Block 4:** Recognize themes and implications - **Block 5:** Make inferences and predictions - **Block 6:** Full contextual understanding

Each level builds on the previous ones!

Key Point: Transformer blocks are the core building blocks that combine attention (communication) and feed-forward (computation) with residual connections (preservation) to build increasingly sophisticated representations.

10. The Complete GPT Architecture

The Full Picture

Now let's put **everything** together into the complete GPT-2 architecture:

```

Text Input: "The cat sat"
↓
Tokenization
↓
Token IDs: [145, 892, 33]
↓
Token Embeddings (384-dim vectors)
↓
+ Position Embeddings (add position info)
↓
Dropout (regularization)
↓

Transformer Block 1
- Layer Norm
- Multi-Head Attention
- Residual Connection

```

- Layer Norm
- Feed-Forward
- Residual Connection

↓

Transformer Block 2
- (same structure)

↓

...
↓

Transformer Block N
- (same structure)

↓

Final Layer Norm

↓

Language Model Head (project to vocab)

↓

Logits: [50257 scores per position]

↓

Softmax (convert to probabilities)

↓

Sample (pick next token)

↓

Output: Next token prediction

Architecture Hyperparameters

Let's use a small GPT as an example:

```
config = {
    'vocab_size': 50257,           # Number of possible tokens
    'embedding_dim': 384,          # Size of embeddings
    'block_size': 256,             # Max sequence length (context window)
    'n_heads': 6,                  # Number of attention heads
    'n_layers': 6,                 # Number of transformer blocks
    'dropout': 0.2                # Dropout probability
}
```

Parameter Count

Let's calculate how many parameters this model has:

```
# Token embeddings: vocab_size × embedding_dim
token_emb = 50257 × 384 = 19,298,688
```

```

# Position embeddings: block_size x embedding_dim
pos_emb = 256 x 384 = 98,304

# Each transformer block has:
# Multi-Head Attention:
#   - Q, K, V projections: 3 x (384 x 64) x 6 heads = 442,368
#   - Output projection: 384 x 384 = 147,456
# Feed-Forward:
#   - Expand: 384 x 1536 = 589,824
#   - Compress: 1536 x 384 = 589,824
# Layer Norms: ~1,536
# Total per block: ~1,771,008

# 6 blocks: 6 x 1,771,008 = 10,626,048

# Final layer norm: 768
# LM head: 384 x 50257 = 19,298,688

# TOTAL: ~49 million parameters!

```

The Three Phases

Phase 1: Encoding (Embeddings)

```

def _embed_tokens(self, idx):
    T = idx.size(1)
    # Convert tokens to vectors
    tok_emb = self.token_embedding(idx)
    # Add position information
    pos_emb = self.position_embedding(torch.arange(T))
    # Combine and regularize
    return self.dropout(tok_emb + pos_emb)

```

What happens: - Token IDs → 384-dim vectors - Each vector knows what it is (content) and where it is (position)

Phase 2: Transformation (Blocks)

```

def _apply_blocks(self, x):
    for block in self.blocks:
        x = block(x)  # Apply attention + feed-forward
    return x

```

What happens: - 6 transformer blocks refine the representations - Each block adds understanding through attention and computation - Information flows through the residual stream

Phase 3: Prediction (Output)

```

# Normalize
x = self.ln_f(x)

```

```

# Project to vocabulary
logits = self.head(x)  # (B, T, vocab_size)

# For each position, we now have 50,257 scores
# Higher score = more likely to be the next token

```

The Complete Forward Pass

```

class GPT2(nn.Module):
    def forward(self, idx, targets=None):
        # Phase 1: Embed tokens with positions
        x = self._embed_tokens(idx)  # (B, T, 384)

        # Phase 2: Apply transformer blocks
        x = self._apply_blocks(x)  # (B, T, 384)

        # Phase 3: Predict next tokens
        x = self.ln_f(x)  # Final normalization
        logits = self.head(x)  # (B, T, 50257)

        # If training, compute loss
        loss = None
        if targets is not None:
            loss = self._compute_loss(logits, targets)

    return logits, loss

```

Example Flow with Real Numbers

```

# Input: "The cat sat" → [145, 892, 33]
idx = torch.tensor([[145, 892, 33]])  # (1, 3)

# Step 1: Token Embeddings
tok_emb = token_embedding(idx)  # (1, 3, 384)
# [[145] → [0.23, -0.41, 0.56, ...], (384 numbers)]
# [892] → [0.12, 0.87, -0.34, ...], (384 numbers)
# [33] → [-0.56, 0.23, 0.91, ...]] (384 numbers)

# Step 2: Add Positions
pos_emb = position_embedding([0, 1, 2])  # (3, 384)
x = tok_emb + pos_emb  # (1, 3, 384)

# Step 3: Transform through 6 blocks
x = block1(x)  # (1, 3, 384)
x = block2(x)  # (1, 3, 384)
x = block3(x)  # (1, 3, 384)
x = block4(x)  # (1, 3, 384)

```

```

x = block5(x)  # (1, 3, 384)
x = block6(x)  # (1, 3, 384)

# Step 4: Final prediction
x = layer_norm(x)  # (1, 3, 384)
logits = lm_head(x)  # (1, 3, 50257)

# For each of the 3 positions, we have 50,257 scores!
# Position 0 ("The"): predicts next token after "The"
# Position 1 ("cat"): predicts next token after "The cat"
# Position 2 ("sat"): predicts next token after "The cat sat"

# Highest scoring token at position 2 might be "on" or "down"

```

Why This Architecture Works

1. **Embeddings:** Transform discrete tokens into continuous vectors that can be processed
2. **Position Info:** Let the model understand sequence order
3. **Attention:** Let each position gather context from previous positions
4. **Feed-Forward:** Process and transform the contextualized information
5. **Residual Connections:** Preserve information flow through many layers
6. **Layer Norms:** Stabilize training
7. **Stacking:** Build increasingly abstract representations
8. **Final Projection:** Convert representations back to vocabulary space

Key Point: The complete GPT architecture is a carefully designed pipeline that transforms text into embeddings, refines those embeddings through multiple transformer blocks, and produces probability distributions over the next token.

11. Training: Teaching the Model

What is Training?

Training is the process of adjusting the model's ~49 million parameters so it gets better at predicting the next token.

The Training Loop

```

for step in range(max_steps):
    # 1. Get a batch of data
    x, y = get_batch(train_data, block_size, batch_size)

    # 2. Forward pass: make predictions
    logits, loss = model(x, y)

    # 3. Backward pass: compute gradients
    loss.backward()

```

```
# 4. Update parameters
optimizer.step()
optimizer.zero_grad()
```

Let's break down each step!

Step 1: Get Training Data

```
def get_batch(data, block_size, batch_size):
    # Pick random starting positions
    ix = torch.randint(len(data) - block_size, (batch_size,))

    # Extract sequences
    x = torch.stack([data[i:i+block_size] for i in ix])
    y = torch.stack([data[i+1:i+block_size+1] for i in ix])

    return x, y
```

Example:

Data: "The cat sat on the mat and slept"
 Block size: 4

Sample 1:

x: "The cat sat on" → Predict each next token
 y: "cat sat on the" → Target for each position

Sample 2:

x: "on the mat and"
 y: "the mat and slept"

The model learns: given x, predict y!

Step 2: Forward Pass & Loss Computation

```
logits, loss = model(x, targets=y)
```

What happens:

```
# x: "The cat sat on" → [145, 892, 33, 670]
# y: "cat sat on the" → [892, 33, 670, 145]

# Model predicts:
logits = model(x) # (batch, 4, 50257)

# For position 0: model predicts what comes after "The"
# Target: "cat" (token 892)
# Model's scores: [0.01, 0.02, ..., 0.95, ..., 0.01]
#                                     ↑
#                                     token 892
# Loss: How wrong was this prediction?
```

```

# Cross-entropy loss:
loss = -log(probability_of_correct_token)

# If model was very confident (prob=0.95): loss = -log(0.95) = 0.05 (good!)
# If model was unsure (prob=0.1): loss = -log(0.1) = 2.3 (bad!)

```

Cross-Entropy Loss measures: “How surprised are you that the correct answer is X?” - Low surprise (high probability) = low loss = good! - High surprise (low probability) = high loss = bad!

Step 3: Backward Pass (Backpropagation)

```
loss.backward()
```

This is where the magic happens! PyTorch automatically computes **gradients**:

What's a gradient? - Tells each parameter: “If you increase, does loss go up or down?” - Tells us: “How should we adjust this parameter to reduce loss?”

```

# For each parameter (49 million of them!):
# gradient = loss / parameter

```

```

# Example:
weight = 0.5 (current value)
gradient = 0.3 (tells us: increasing weight increases loss)

# So we should DECREASE this weight!

```

Step 4: Parameter Update (Optimization)

```
optimizer.step() # Update parameters using gradients
optimizer.zero_grad() # Reset gradients for next iteration
```

Adam Optimizer (most common):

```

# For each parameter:
new_value = old_value - learning_rate * gradient

```

```

# Example:
old_weight = 0.5
gradient = 0.3 (loss increases with weight)
learning_rate = 0.001

```

```
new_weight = 0.5 - 0.001 * 0.3 = 0.4997
```

```
# Weight decreased slightly, which should reduce loss!
```

The Learning Rate

Critical hyperparameter!

```

# Too large (lr = 0.1):
weight = 0.5 - 0.1 * 0.3 = 0.47 # Big jumps, might overshoot!

# Too small (lr = 0.0001):
weight = 0.5 - 0.0001 * 0.3 = 0.49997 # Tiny steps, very slow!

# Just right (lr = 0.001):
weight = 0.5 - 0.001 * 0.3 = 0.4997 # Steady progress!

```

Gradient Clipping

```
torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
```

Why? Sometimes gradients get HUGE (exploding gradients):

```

# Without clipping:
gradient = 10000 (huge!)
update = -0.001 * 10000 = -10 (massive change, breaks training!)

# With clipping (max=1.0):
if gradient > 1.0:
    gradient = 1.0
update = -0.001 * 1.0 = -0.001 (reasonable change)

```

The Complete Training Step

```

def _train_step(model, optimizer, train_data, block_size, batch_size, device):
    # 1. Sample batch
    xb, yb = get_batch(train_data, block_size, batch_size, device)

    # 2. Forward pass
    logits, loss = model(xb, yb)

    # 3. Zero gradients from previous step
    optimizer.zero_grad()

    # 4. Backward pass
    loss.backward()

    # 5. Clip gradients (prevent explosion)
    torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)

    # 6. Update parameters
    optimizer.step()

    # 7. Return loss for monitoring
    return loss.item()

```

Evaluation

We also evaluate on a **validation set** (data the model hasn't seen during training):

```
def _eval_model(model, val_data, block_size, batch_size, device, eval_iters):
    model.eval() # Turn off dropout
    losses = []

    with torch.no_grad(): # Don't compute gradients
        for _ in range(eval_iters):
            xv, yv = get_batch(val_data, block_size, batch_size, device)
            _, loss = model(xv, yv)
            losses.append(loss.item())

    model.train() # Turn dropout back on
    return sum(losses) / len(losses)
```

Why evaluate? - Check if model is actually learning (validation loss should decrease) - Detect overfitting (validation loss increases while training loss decreases)

Training Progress

Step 0	Train Loss:	10.83	Val Loss:	10.85	(random initialization)
Step 100	Train Loss:	6.45	Val Loss:	6.52	(learning basic patterns)
Step 500	Train Loss:	4.23	Val Loss:	4.31	(learning syntax)
Step 1000	Train Loss:	3.15	Val Loss:	3.28	(learning semantics)
Step 5000	Train Loss:	2.31	Val Loss:	2.45	(generating coherent text)
Step 10000	Train Loss:	1.89	Val Loss:	2.12	(good performance!)

Checkpointing

Save the best model:

```
def save_checkpoint(model, optimizer, step, val_loss, filepath):
    torch.save({
        'step': step,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'val_loss': val_loss,
    }, filepath)
```

Why? - Training might crash (save progress!) - Want to use the best model (not the last) - Can resume training later

The Big Picture

```
Start: Random parameters (model outputs gibberish)
      ↓
Train for 100 steps (model learns common words)
      ↓
Train for 1000 steps (model learns grammar)
```

```

↓
Train for 10000 steps (model writes coherent sentences)
↓
End: Trained model (can generate human-like text!)

```

Key Point: Training adjusts millions of parameters through gradient descent, gradually teaching the model to predict the next token by showing it countless examples and adjusting parameters to minimize prediction errors.

12. Generation: Creating New Text

From Training to Generation

Training: Learn to predict next tokens from examples **Generation:** Use those learned patterns to create new text!

The Generation Process

```

def generate(model, start_tokens, max_new_tokens, temperature=1.0):
    """
    Start with some tokens, generate new ones!

    Args:
        start_tokens: e.g., "The cat" → [145, 892]
        max_new_tokens: How many new tokens to generate
        temperature: Controls randomness
    """
    current_tokens = start_tokens

    for _ in range(max_new_tokens):
        # 1. Get model predictions
        logits, _ = model(current_tokens)

        # 2. Focus on last position (next token)
        next_token_logits = logits[:, -1, :]

        # 3. Apply temperature
        next_token_logits = next_token_logits / temperature

        # 4. Convert to probabilities
        probs = F.softmax(next_token_logits, dim=-1)

        # 5. Sample next token
        next_token = torch.multinomial(probs, num_samples=1)

        # 6. Append to sequence
        current_tokens = torch.cat([current_tokens, next_token], dim=1)

```

```
    return current_tokens
```

Let's break this down!

Step-by-Step Generation Example

Initial Context

Input: "The cat sat on the"

Tokens: [145, 892, 33, 670, 145]

Iteration 1: Generate Token 6

```
# 1. Forward pass
logits, _ = model([145, 892, 33, 670, 145])
# logits shape: (1, 5, 50257)
#           ↑ ↑ ↑
#       batch, seq_len, vocab_size

# 2. Get predictions for position after "the"
next_logits = logits[:, -1, :]
# These are raw scores for each possible next token

# Example logits (just first 10 tokens):
# token 0 ("!"): -2.3
# token 1 (.): -1.5
# token 2 ("the"): -3.1
# token 3 ("mat"): 2.8 ← High score!
# token 4 ("dog"): 1.2
# ...

# 3. Apply temperature (we'll explain this next!)
next_logits = next_logits / 1.0 # temperature = 1.0

# 4. Convert to probabilities with softmax
probs = softmax(next_logits)
# token 3 ("mat"): 0.45 ← 45% probability
# token 4 ("dog"): 0.12 ← 12% probability
# token 8 ("floor"): 0.08 ← 8% probability
# ... all sum to 1.0

# 5. Sample from probability distribution
next_token = sample(probs) # Let's say we get token 3 ("mat")

# 6. Append to sequence
tokens = [145, 892, 33, 670, 145, 3]
text = "The cat sat on the mat"
```

Iteration 2: Generate Token 7

```
# Now we have: "The cat sat on the mat"  
# Model predicts: What comes after "mat"?  
  
logits, _ = model([145, 892, 33, 670, 145, 3])  
next_logits = logits[:, -1, :] # Position after "mat"  
  
# Possible predictions:  
# "." : 0.35 (35% - end the sentence)  
# "and": 0.25 (25% - continue the sentence)  
# "," : 0.18 (18% - add a clause)  
  
# Sample: Let's say we get "."  
tokens = [145, 892, 33, 670, 145, 3, 13]  
text = "The cat sat on the mat."
```

Temperature: Controlling Randomness

Temperature is crucial! It controls how “creative” vs “deterministic” the model is.

Low Temperature (temperature < 1.0) - Focused, Deterministic

```
# Original logits:  
logits = [1.0, 2.0, 3.0, 1.5]  
  
# With temperature = 0.5 (divide by 0.5 = multiply by 2)  
scaled = [2.0, 4.0, 6.0, 3.0] # Differences amplified!  
  
# After softmax:  
probs = [0.01, 0.12, 0.86, 0.01]  
#           ↑  
#       Model is very confident about this choice!
```

Effect: Model picks the most likely token almost every time - **Pros:** More coherent, grammatically correct - **Cons:** Boring, repetitive (“the the the...”)

High Temperature (temperature > 1.0) - Creative, Random

```
# Original logits:  
logits = [1.0, 2.0, 3.0, 1.5]  
  
# With temperature = 2.0  
scaled = [0.5, 1.0, 1.5, 0.75] # Differences reduced!  
  
# After softmax:  
probs = [0.19, 0.29, 0.38, 0.14]  
#       More evenly distributed!
```

Effect: Model considers less likely tokens - **Pros:** More creative, varied, interesting - **Cons:** Can be incoherent or make mistakes

Temperature = 1.0 - Balanced

```
# Original logits (unchanged)
probs = softmax(logits) # Normal sampling
```

Effect: Sample exactly as the model learned - Good default choice!

Special Case: Temperature = 0 (Greedy Decoding)

```
# Don't sample - just pick the highest probability token
next_token = argmax(probs)
```

Effect: Always pick the most likely token - Most deterministic - Useful for reproducible output

Examples with Different Temperatures

Prompt: "Once upon a time, there was a"

Temperature = 0.5 (Conservative):

"Once upon a time, there was a little girl who lived in a small village with her mother and father."

(Predictable, grammatically perfect)

Temperature = 1.0 (Balanced):

"Once upon a time, there was a young wizard who discovered a mysterious book in his grandfather's attic."

(Coherent but more interesting)

Temperature = 1.5 (Creative):

"Once upon a time, there was a talking fish who dreamed of becoming a ballet dancer on the moon."

(Creative but might get weird!)

Temperature = 2.0 (Very Random):

"Once upon a time, there was a purple mathematics equation that sang opera to the refrigerator's grandmother."

(Too random, incoherent)

Context Window Management

GPT has a **limited context window** (e.g., 256 tokens). What if our generated text exceeds this?

```
def _generate_step(self, idx, temperature):
    # Crop to last block_size tokens
    idx_cond = idx[:, -self.block_size:]
```

```

# Forward pass with cropped context
logits, _ = self(idx_cond)

# Rest of generation...

```

Example:

Context window: 256 tokens
Generated so far: 300 tokens

Use only: tokens[300-256:300] = last 256 tokens
Model doesn't see: the first 44 tokens

This is why GPT can "forget" earlier context in long conversations!

The Complete Generation Function

```

def generate(self, idx, max_new_tokens, temperature=1.0):
    """Generate new tokens autoregressively."""
    for _ in range(max_new_tokens):
        # Crop context if needed
        idx_cond = idx[:, -self.block_size:]

        # Get predictions
        logits, _ = self(idx_cond)

        # Focus on last position
        logits = logits[:, -1, :] / temperature

        # Convert to probabilities
        probs = F.softmax(logits, dim=-1)

        # Sample next token
        next_token = torch.multinomial(probs, num_samples=1)

        # Append
        idx = torch.cat([idx, next_token], dim=1)

    return idx

```

Beam Search (Alternative to Sampling)

Instead of sampling one token at a time, we can keep track of multiple possibilities:

Start: "The cat"

Beam 1: "The cat sat" (prob=0.8)
Beam 2: "The cat jumped" (prob=0.6)
Beam 3: "The cat ran" (prob=0.5)

Continue each beam...

Pick the sequence with highest total probability!

Not implemented in GPT-2, but used in some applications.

Key Point: Generation is the reverse of training - instead of predicting the next token from known text, we use the model's learned predictions to create new text autoregressively, with temperature controlling the balance between coherence and creativity.

13. Putting It All Together

The Complete Pipeline

Let's trace a single example through the **entire system**:

Input: “The cat sat on the mat”

Phase 1: Tokenization

```
text = "The cat sat on the mat"
tokens = tokenizer.encode(text)
# [145, 892, 33, 670, 145, 3]
```

Each word/piece gets a unique ID.

Phase 2: Embedding

```
# Token embeddings (what are these tokens?)
tok_emb = token_embedding([145, 892, 33, 670, 145, 3])
# Shape: (6, 384)
# [[0.23, -0.41, 0.56, ...],   ← "The"
# [0.12, 0.87, -0.34, ...],   ← "cat"
# [-0.56, 0.23, 0.91, ...],   ← "sat"
# [0.34, -0.12, 0.45, ...],   ← "on"
# [0.23, -0.41, 0.56, ...],   ← "the" (same as position 0!)
# [0.78, 0.34, -0.23, ...]]   ← "mat"

# Position embeddings (where are these tokens?)
pos_emb = position_embedding([0, 1, 2, 3, 4, 5])
# Shape: (6, 384)
# [[0.11, 0.22, -0.33, ...],   ← position 0
# [0.44, -0.55, 0.66, ...],   ← position 1
# [0.77, 0.88, -0.99, ...],   ← position 2
# [0.12, -0.23, 0.34, ...],   ← position 3
# [0.45, 0.56, -0.67, ...],   ← position 4
# [0.78, -0.89, 0.90, ...]]   ← position 5

# Combined embeddings
x = tok_emb + pos_emb # (6, 384)
# Each position now knows WHAT it is and WHERE it is!
```

Phase 3: Transformer Block 1

```
# Input: (6, 384)

# Sub-layer 1: Multi-Head Attention
x_norm = layer_norm(x)

# Each position looks at previous positions:
# "The" (pos 0): Only sees itself
# "cat" (pos 1): Sees "The" and "cat"
# "sat" (pos 2): Sees "The", "cat", "sat"
# "on" (pos 3): Sees "The", "cat", "sat", "on"
# "the" (pos 4): Sees all previous
# "mat" (pos 5): Sees all previous

attn_out = multi_head_attention(x_norm) # (6, 384)
x = x + attn_out # Residual connection

# Sub-layer 2: Feed-Forward
x_norm = layer_norm(x)
ff_out = feed_forward(x_norm) # (6, 384)
x = x + ff_out # Residual connection

# After Block 1: (6, 384)
# Each position now has context-aware representations!
```

Phase 4: Transformer Blocks 2-6

```
# Same process 5 more times...
x = block2(x) # (6, 384)
x = block3(x) # (6, 384)
x = block4(x) # (6, 384)
x = block5(x) # (6, 384)
x = block6(x) # (6, 384)

# After 6 blocks:
# "mat" has gathered information from:
# - Direct context: "on the"
# - Subject: "cat"
# - Action: "sat"
# - Full sentence structure and meaning
```

Phase 5: Output Projection

```
# Final layer norm
x = layer_norm(x) # (6, 384)

# Project to vocabulary
logits = lm_head(x) # (6, 50257)
```

```
# For each position, we have 50,257 scores (one per token!)
```

Phase 6: Predictions

```
# Position 0: "The" predicts next token
#   Highest scores: "cat" (892), "dog" (1234), "man" (5678)
#   Model predicts: "cat"  Correct!

# Position 1: "The cat" predicts next token
#   Highest scores: "sat" (33), "ran" (789), "jumped" (456)
#   Model predicts: "sat"  Correct!

# Position 2: "The cat sat" predicts next token
#   Highest scores: "on" (670), "down" (234), "still" (567)
#   Model predicts: "on"  Correct!

# Position 3: "The cat sat on" predicts next token
#   Highest scores: "the" (145), "a" (678), "my" (901)
#   Model predicts: "the"  Correct!

# Position 4: "The cat sat on the" predicts next token
#   Highest scores: "mat" (3), "floor" (890), "couch" (567)
#   Model predicts: "mat"  Correct!

# Position 5: "The cat sat on the mat" predicts next token
#   Highest scores: "." (13), "and" (234), "," (456)
#   Model predicts: "." (end of sentence!)
```

Understanding What Each Component Does

Token + Position Embeddings

- **Input:** Discrete token IDs
- **Output:** Dense vectors with content and position info
- **Purpose:** Create a rich, continuous representation

Multi-Head Attention

- **Input:** Contextualized vectors from previous layer
- **Output:** Vectors enriched with information from other positions
- **Purpose:** “Reading comprehension” - understanding relationships

Feed-Forward Networks

- **Input:** Attention-enriched vectors
- **Output:** Transformed vectors with extracted features
- **Purpose:** “Thinking” - processing information at each position

Residual Connections

- **Input/Output:** Same shape
- **Purpose:** “Information highway” - preserve original information

Layer Normalization

- **Input:** Any vector
- **Output:** Normalized vector (mean=0, std=1)
- **Purpose:** Stabilize training, prevent exploding/vanishing values

Final Projection

- **Input:** Final transformer representations (384-dim)
- **Output:** Vocabulary logits (50257-dim)
- **Purpose:** Convert learned representations to token predictions

Training vs Generation

During Training:

```
# We have both input and targets
x = "The cat sat on the"
y = "cat sat on the mat" # Shifted by 1

# Forward pass
logits, loss = model(x, y)

# Compute how wrong predictions are
loss.backward() # Compute gradients
optimizer.step() # Update parameters

# Model gets slightly better at predicting!
```

During Generation:

```
# We only have input (no targets)
x = "The cat sat on the"

# Forward pass
logits, _ = model(x, targets=None) # No loss

# Sample next token
next_token = sample(logits[:, -1, :]) # "mat"

# Append and repeat
x = "The cat sat on the mat"
next_token = sample(logits[:, -1, :]) # "."

# Result: "The cat sat on the mat."
```

Why GPT Works So Well

1. **Self-Attention:** Can look at any previous position (long-range dependencies!)
2. **Multi-Head:** Learns multiple types of relationships simultaneously
3. **Depth:** 6+ layers build increasingly abstract representations
4. **Residual Stream:** Information flows smoothly through many layers
5. **Scale:** 49M+ parameters can memorize patterns and facts
6. **Autoregressive:** Simple training objective (predict next token)
7. **Massive Data:** Trained on billions of tokens from the internet

What Makes It “Intelligent”?

GPT doesn't truly “understand” - it recognizes patterns:

Training data includes:

"The cat sat on the mat"
"The dog sat on the floor"
"The bird sat on the tree"

GPT learns:

- [animal] + "sat on the" + [location]
- Verbs agree with subjects
- Prepositions connect actions to places
- Sentence structure patterns

When generating:

"The cat sat on the ___"
→ Likely completions: mat, floor, chair, couch
→ Unlikely: sky, galaxy, thought (doesn't fit pattern!)

Limitations

1. **Context Window:** Can only see last N tokens (256, 2048, etc.)
2. **No Real Understanding:** Pattern matching, not reasoning
3. **Training Data:** Only knows what it was trained on
4. **No Memory:** Each generation is independent (unless you include conversation history)
5. **Computational Cost:** Requires powerful GPUs for training
6. **Factual Errors:** Can confidently generate wrong information

But It's Still Amazing!

From simple principles: - “Predict the next token” - Self-attention - Feed-forward layers - Residual connections

We get a model that can: - Write coherent essays - Translate languages - Answer questions - Write code - Have conversations

All from learning to predict the next token!

Conclusion

What We've Learned

You now understand every component of GPT:

1. **Tokenization:** Text → Numbers
2. **Token Embeddings:** Numbers → Vectors
3. **Position Embeddings:** Adding position information
4. **Residual Connections:** Information superhighway
5. **Self-Attention:** Looking at context
6. **Multi-Head Attention:** Multiple perspectives
7. **Feed-Forward Networks:** Processing information
8. **Transformer Blocks:** Combining it all
9. **Complete Architecture:** The full pipeline
10. **Training:** Teaching the model
11. **Generation:** Creating new text

The Core Insight

GPT is “just” predicting the next token, but by doing this billions of times on massive amounts of text, it learns:
- Grammar and syntax - Facts and knowledge - Reasoning patterns
- Writing styles - And much more!

Next Steps

1. **Run the code:** Train your own GPT-2!
2. **Experiment:** Try different hyperparameters
3. **Visualize:** Plot attention weights to see what the model focuses on
4. **Scale up:** Train a bigger model (more layers, larger n_embd)
5. **Fine-tune:** Train on specific data (code, poems, dialogue)

Key Takeaway

Understanding GPT demystifies AI:
- It's not magic - It's clever mathematics and engineering - It's accessible (you can build one!) - But it's also incredibly powerful when scaled up

You now know how ChatGPT-style models work under the hood!

Appendix: Quick Reference

Dimensions Cheat Sheet

```
vocab_size = 50257          # Number of tokens
embedding_dim = 384         # Size of vectors
block_size = 256            # Max sequence length
n_heads = 6                 # Number of attention heads
head_size = 64               # embedding_dim // n_heads
n_layers = 6                # Number of transformer blocks
batch_size = 4               # Number of sequences processed together
```

```

# Shapes through the network:
Input tokens: (batch_size, seq_len)
→ (4, 256)

After embeddings: (batch_size, seq_len, embedding_dim)
→ (4, 256, 384)

After attention head: (batch_size, seq_len, head_size)
→ (4, 256, 64)

After multi-head attention: (batch_size, seq_len, embedding_dim)
→ (4, 256, 384)

After feed-forward: (batch_size, seq_len, embedding_dim)
→ (4, 256, 384)

After transformer block: (batch_size, seq_len, embedding_dim)
→ (4, 256, 384)

Final logits: (batch_size, seq_len, vocab_size)
→ (4, 256, 50257)

```

Hyperparameter Guidelines

Parameter	Tiny	Small	Medium	Large
embedding_dim	64	384	768	1024
n_layers	2	6	12	24
n_heads	2	6	12	16
block_size	64	256	1024	2048
Parameters	~1M	~49M	~125M	~350M

Important Equations

Attention:

$$\text{Attention}(Q, K, V) = \text{softmax}(Q @ K^T / \sqrt{d_k}) @ V$$

Feed-Forward:

$$\text{FFN}(x) = \text{ReLU}(x @ W_1 + b_1) @ W_2 + b_2$$

Residual Connection:

$$\text{output} = \text{input} + \text{layer}(\text{input})$$

Layer Norm:

$$\text{LayerNorm}(x) = \frac{x - \text{mean}}{\text{std}} + \text{scale}$$

Cross-Entropy Loss:

```
Loss = -log(P(correct_token))
```

Key Concepts Summary

- **Autoregressive:** Predict one token at a time, based on all previous
 - **Causal Masking:** Can't look at future tokens
 - **Temperature:** Controls randomness in generation
 - **Gradient Descent:** How we train neural networks
 - **Backpropagation:** How we compute gradients efficiently
 - **Attention:** Mechanism for focusing on relevant information
 - **Residual Stream:** Information highway through the network
-

Happy Learning!

For code implementation, see: `micro_gpt.py`

For comprehensive tests, see: `test_micro_gpt.py`

For running examples, see: `example.py`